

Sreenivasulu Ummadi

ummadi.sreenivasulu.developer@gmail.com | linkedin.com/in/sreenivasuluummadi/ | github.com/sreenu-dev |
+91-8639912976

TECHNICAL SKILLS

Languages: Python, Java, JavaScript, SQL, TypeScript, c

Frameworks: React, Node.js, Angular, ASP.NET Web API, xUnit, NUnit

Developer Tools: Git, Gitlab, Kubernetes, Azure, PostgreSQL, Redis, Jira, Kafka

Architecture & Concepts: Microservices, Latency Optimization, Micro Frontend Architecture, System Design, Data Structure, Algorithms, OAuth2, JWT, IAM, Data Encryption(AES/TLS)

EXPERIENCE

Deloitte

Consultant

Hyderabad

June 2025 – Present

- Engineered v2 backend API endpoints, reducing the system error ratio by 70% and accelerating response velocity by designing a highly structured, optimized response schema.
- Modernized legacy data access layers by migrating codebase modules to the latest .NET scaffolding architecture, allowing direct database object-mapping via Entity Framework and eliminating redundant boilerplate queries.
- Optimized database performance and query throughput by tuning legacy SQL stored procedures and refactoring the Entity Framework data execution context to handle complex relational outputs efficiently, improving query efficiency by 25%.
- Enhanced system observability and production troubleshooting by replacing a traditional logging mechanism with a custom structured logger integrated with OpenTelemetry, capable of parsing granular success/error contextual data payloads.
- Integrated a highly scalable asset-upload feature into the core ecosystem, supporting high-throughput file transfers while maintaining strict backward compatibility to ensure zero downtime for existing user workflows.
- Decoupled a monolithic frontend application by migrating to a Micro-Frontend architecture, significantly eliminating cross-team dependency bottlenecks and reducing local build deployment overhead.
- Standardized user interface consistency across distributed application modules by introducing centralized, global styling packages, reducing design regression friction by ensuring uniform component behaviors.
- Accelerated pipeline automation and deployment velocity by authoring custom Kubernetes (K8s) deployment configuration manifests to manage containerized orchestration across the CI/CD pipeline environment.
- Boosted testing reliability and accelerated code delivery cycles by engineering mock payloads and establishing automated unit-testing workflows across the entire application stack.
- Reduced developer onboarding overhead and streamlined team scaling by architecting comprehensive Confluence technical playbooks, documentation architectures, and system setup guidelines to eliminate manual knowledge transfer (KT) dependencies.
- Accelerated code delivery and scaled test reliability by orchestrating autonomous Agentic AI agents to automatically synthesize comprehensive unit test suites and complex edge-case mock payloads, reducing manual test authorship overhead by 40%.
- Hardened codebase security and minimized architectural regressions by implementing an LLM-driven agentic review pipeline that flags structural code risks, analyzes vulnerability vectors, and outputs automated mitigation scripts directly within the CI/CD pull request workflow.
- Boosted engineering velocity and optimized the software development lifecycle (SDLC) by deploying multi-agent AI workflows to automate structural code generation and repetitive boilerplate scaffolding, improving feature time-to-market by 25%.
- Proactively mitigated deployment risks and production downtime by building an intelligent agentic monitoring assistant that automatically parses system logs, identifies anomalous behavior patterns, and suggests real-time remediation strategies for the engineering team.

Accenture

Full Stack Engineering Senior Analyst

Hyderabad

March 2022 – June 2025

- Optimized application throughput and system scalability by 20% by designing and implementing complex, high-performance features aligned with detailed technical specifications.
- Maintained a 99.9% production uptime rating by debugging and resolving critical, cross-stack system issues across development, testing, and live environments.

- Reduced post-release software regression rates by 15% by engineering robust technical solutions for new product features and critical system bug fixes.
- Hardened codebase reliability and achieved zero regression failures by architecting comprehensive front-end unit test suites utilizing Jasmine and Karma frameworks.
- Decreased future technical debt and development lifecycle overhead by 15% through the continuous delivery of clean, reusable, and scalable code adhering to industry best practices.
- Streamlined cross-functional alignment and cut feature design phases by 20% by leading project steering meetings, directing the Request-for-Comment (RFC) process, and providing technical leadership.
- Accelerated engineering onboarding timelines by 30% by actively supporting team initiatives, organizing system architectures, and mentoring junior team members.
- Expanded engineering team capability by 25% while maintaining top-tier quality benchmarks by mentoring software engineers and driving the technical interview process for top talent.
- Boosted sprint velocity and on-time feature delivery by 18% by continuously monitoring core team performance metrics and proactively mitigating development bottlenecks.
- Accelerated product roadmap velocity by 25% by engineering 100+ end-to-end user stories and high-scalability features while maintaining modular design principles.
- Enhanced core application stability and minimized production downtime by debugging and resolving 300+ high-priority, systemic defects across distributed environments.

Accenture

Bengaluru

Full Stack Engineering Analyst

September 2020 – March 2022

- Boosted application maintainability and reduced cross-device rendering latency by 20% by designing and building modular, scalable front-end architectures using Angular (v8+) and TypeScript.
- Improved mobile user retention by 15% and enhanced cross-device visual appeal by architecting and maintaining responsive, user-facing features utilizing HTML5, CSS3, SCSS, and JavaScript.
- Decreased front-end feature development cycle times by 25% by implementing and optimizing reusable, component-based design patterns to maximize codebase modularity.
- Reduced post-release architectural technical debt by 15% by applying strict Object-Oriented Programming (OOP) principles to engineer robust, high-quality front-end codebases.
- Accelerated page load speeds by 35% and optimized cross-device responsiveness by refactoring asset delivery pipelines and eliminating client-side rendering bottlenecks.
- Reduced end-to-end data synchronization latency by 30% by partnering with cross-functional back-end engineering teams to seamlessly integrate high-throughput RESTful APIs.
- Elevated application stability and improved user experience scores by 20% by troubleshooting and resolving high-priority front-end performance bottlenecks and critical production bugs.
- Hardened codebase reliability and achieved a 95% automated test coverage rate by architecting comprehensive front-end unit test suites utilizing Jasmine and Karma frameworks to eliminate production regressions.
- Increased sprint velocity by 15% and elevated team code quality standards by actively driving agile practices, including detailed sprint planning, thorough peer code reviews, and daily stand-ups.
- Achieved 100% WCAG compliance and expanded the accessible user base by auditing and refactoring legacy UI components to adhere strictly to semantic web standards and inclusivity guidelines.
- Accelerated feature time-to-market by 18% by partnering closely with cross-functional product, design, and business teams to align technical front-end solutions with core product requirements.

PROJECTS

Distributed Event-Driven Ledger | *.NET Core, Apache Kafka, Redis, PostgreSQL, Kubernetes* Jan 2026 – Mar 20

- Architected a fault-tolerant financial ledger microservice using .NET Core and Apache Kafka, implementing CQRS and Event Sourcing patterns to seamlessly ingest 20,000+ simulated concurrent transactions per second without data loss.
- Eliminated transactional race conditions and double-spending vulnerabilities across distributed nodes by engineering a Redis-driven distributed locking mechanism.
- Optimized read-path latency by 40% by separating command and query pipelines, utilizing Kafka consumers to asynchronously update a read-optimized PostgreSQL database instance.
- Guaranteed 99.99% data consistency and auditability by persisting immutable event streams in Kafka, allowing for deterministic state reconstruction during simulated system failures.

Enterprise GenAI Compliance Guardrail Engine | *Python, LangGraph, OpenTelemetry, pgvector, Azure* March 20

- Engineered an autonomous financial compliance guardrail pipeline using Python and LangGraph to parse unstructured audit logs, yielding a 94% accuracy rate in automated regulatory intent classification.
- Minimized runtime inferencing costs and optimized latency by 30% by deploying custom Redis semantic caching layers and utilizing OpenTelemetry to track baseline performance tokens.
- Implemented an asynchronous validation layer using task queues to run parallel multi-agent evaluation chains, reducing end-to-end compliance report generation time by 35%.
- Mitigated prompt injection risks and model hallucinations by building structural validation middleware that programmatically checks LLM outputs against strict banking regulatory schemas.

Resilient API Gateway | *Node.js, Express, Redis, Docker, Kubernetes*

May 2026 – Present

- Developed a high-availability API Gateway utilizing the Strangler and Aggregation patterns to safely decouple monolithic endpoints into containerized Kubernetes microservices.
- Mitigated cascading backend system degradations under a simulated load of 15,000 requests per minute by engineering distributed Circuit Breaker algorithms and token-bucket rate limiters.
- Reduced microservice discovery overhead and cross-pod routing latency by 25% by configuring custom reverse-proxy matching rules and keeping persistent connection pools.
- Enhanced system fault-tolerance by implementing a graceful fallback mechanism that serves stale cached data from Redis when core downstream dependencies fail.

EDUCATION

Karunya University

Bachelor of Technology in Computer Science and Engineering

Coimbatore, Tamilnadu

2020

CERTIFICATION

- AI-900 Azure AI Fundamentals
- AZ-900 Azure Fundamentals

ACCOMPLISHMENT

- ACE'22 Preformer of the year Award, Accenture
- ACE'23 Extra Mile Award for offshore Individual, Accenture
- FY26 Applause Award, Deloitte